

Advanced Deep Learning

Section 3:

Natural Language Processing



Rényi Alfréd Matematikai Kutatóintézet



ELTE EÖTVÖS LORÁND
TUDOMÁNYEGYETEM

Schedule

Lectures on NLP:

- **Lecture 1: Introduction and Foundations**
 - Characteristics of the domain
 - Classical methods
 - Character encodings
 - Tokenization
 - Embeddings
- **Lecture 2: Language Models and Language Modeling**
 - Objective functions
 - Sequential modeling
 - Decoding strategies
 - Models: Transformers (BERT, GPT)
 - Training: pre-training, fine-tuning
 - Evaluation
- **Lecture 3: Large Language Models (LLMs)**
 - Emergent properties
 - Scaling laws
 - GPT-series
 - Instruction tuning
 - Reinforcement Learning with Human Feedback (RLHF)
- **Lecture 4: Research**
 - Prompt engineering
 - Multimodality: CLIP
 - Problems: hallucination
 - Retrieval Augmented Generation (RAG)
 - Security

Large Language Models (LLMs)

Lecture 3

What is this lecture about?

This lecture tries to answer the following questions:

- What are Large Language Models (LLMs)?
- Why do we increase the size of the Language Models?
- What makes Large Language Models such successful?
 - What are Emergent Abilities?
- What makes a Next Word Prediction model to be able to function in an Application?

Resource

The lecture content has been developed mainly from the paper:

- [A Survey of Large Language Models](#)

Introduction: LLMs

Large Language Models (LLMs):

- Deep Learning models:
 - designed to understand, generate, and manipulate human language in textual form
- Model Architecture:
 - typically Transformer-based
- Model size:
 - ranging from hundreds of millions to billions of parameters
- Capabilities:
 - able to perform a wide range of language-based tasks without task-specific tuning
- Advantages:
 - Versatility: performing multiple tasks
 - Efficiency: generating language with minimal human intervention
 - Real-World Applications: Customer Service; Content Generation, ...

Scaling Laws

Recent trends: Continued Scaling:

- effort to build even larger and more powerful models continue

Scaling Laws:

- a quantitative approach to characterizing the scaling effect
- scaling can largely improve the model capacity:
 - Model size: number of parameters
 - Dataset size: number of training examples
 - Total compute: amount of training compute (time, iteration)

Scaling Laws

KM scaling law:

$$L(N) = \left(\frac{N_c}{N} \right)^{\alpha_N}, \quad \alpha_N \sim 0.076, N_c \sim 8.8 \times 10^{13}$$

$$L(D) = \left(\frac{D_c}{D} \right)^{\alpha_D}, \quad \alpha_D \sim 0.095, D_c \sim 5.4 \times 10^{13}$$

$$L(C) = \left(\frac{C_c}{C} \right)^{\alpha_C}, \quad \alpha_C \sim 0.050, C_c \sim 3.1 \times 10^8$$

Where:

- $L(\bullet)$: the cross entropy loss in nats
- N : model size
- D : dataset size
- C : amount of training compute

Scaling Laws

Chinchilla scaling law:

$$L(N, D) = E + \frac{A}{N^\alpha} + \frac{B}{D^\beta}$$

$$E = 1.69, A = 406.4, B = 410.7, \alpha = 0.34, \beta = 0.28$$

Where:

- $L(\bullet)$: the cross entropy loss in nats
- N : model size
- D : dataset size

Scaling: performance

More text → Lower cross-entropy

More parameters → Lower cross-entropy


Lower cross-entropy → Better performance

- correlating well with other evaluation metrics

Currently scaling does not show signs of topping out

Diagram of sizes

Some of the Popular Large Language Models (LLMs)



Model parameter sizes

Name	Release date ^[a]	Developer	Number of parameters (billion) ^[b]	Corpus size	Training cost (petaFLOP-day)	License ^[c]	Notes
GPT-1	June 2018	OpenAI	0.117		1 ^[121]	MIT ^[122]	First GPT model, decoder-only transformer. Trained for 30 days on 8 P600 GPUs.
BERT	October 2018	Google	0.340 ^[123]	3.3 billion words ^[123]	9 ^[124]	Apache 2.0 ^[125]	An early and influential language model, ^[7] but encoder-only and thus not built to be prompted or generative ^[126]
GPT-2	February 2019	OpenAI	1.5 ^[134]	40GB ^[135] (~10 billion tokens) ^[136]		MIT ^[137]	general-purpose model based on transformer architecture
XLNet	June 2019	Google	~0.340 ^[130]	33 billion words		Apache 2.0 ^[131]	An alternative to BERT; designed as encoder-only ^{[132][133]}
GPT-3	May 2020	OpenAI	175 ^[39]	300 billion tokens ^[136]	3640 ^[138]	proprietary	A fine-tuned variant of GPT-3, termed GPT-3.5, was made available to the public through a web interface called ChatGPT in 2022. ^[139]

Emergent Abilities

Emergent Abilities:

- Qualitative changes in the model's behaviour
 - not trained explicitly
- **Emergent**: an ability is emergent if it is not present in smaller models but is present in larger models
 - arise from the quantitative increase in the scale of the Language Model
- the performance rises significantly when above random when the scale reaches a certain level
- emergent abilities cannot be predicted simply by extrapolating the performance of smaller models
 - the abilities appear suddenly and unpredictable, as the model reaches a critical threshold of scale, rather than improving steadily and predictably
- **Examples:**
 - In-Context Learning
 - Instruction following
 - Chain-of-Thought Prompting

Emergent Abilities: In-Context Learning

In-Context Learning (ICL):

- the ability of a Language Model to learn from the context provided in the input itself, without additional training or external data
 - a natural language instruction and / or
 - several task demonstrations
- introduced by GPT-3
- relying on the examples included in the input prompt to guide its responses

Zero-, One-, Few-shot Learning:

- training the model (updating the parameters) with one / few examples

Demo:

- ChatGPT: eszperente, arithmetic tasks

Emergent Abilities: Instruction following

Instruction following:

- fine-tuning the Language Model with datasets formatted via natural language descriptions (instruction tuning)
- LLMs perform well on unseen tasks that are also described in the form of instructions

Emergent Abilities: Chain-of-Thought Prompting

Chain of Thought (CoT) Prompting:

- a prompt engineering technique
- aims to improve language models' performance on tasks requiring logic, calculation and decision-making
- by structuring the input prompt in a way that mimics human reasoning
- step-by-step reasoning
- Assumption: obtained by training on code

Training phases

Training phases / approaches:

- Pre-Training
- Adaptation / Fine-Tuning:
 - Instruction Tuning
 - Alignment Tuning
 - Reinforcement Learning with Human Feedback (RLHF)
- Prompting

Pre-training

Pre-training:

- optimizing for completion
 - ML task: Language Modeling
- training data:
 - large quantity
 - low quality
 - Scale: 1 trillion tokens (~ 15 million books)
- Examples:
 - GPT-x, LLaMa

Adaptation

Adaptation of LLMs:

- after pre-training
 - LLMs acquired the general abilities
- adapting LLM's abilities to specific goals
- 2 kinds of adaptations:
 - Instruction Tuning
 - Alignment Tuning

Adaptation: Instruction Tuning

Instruction Tuning:

- fine-tuning pre-trained LLMs on a collection of formatted instances in the form of natural language
- Supervised Fine-Tuning
- Dataset:
 - instruction-formatted data
- Supervised training:
 - sequence-to-sequence loss
 - Input:
 - task description (instruction)
 - additional demonstrations (optional)
 - Output:
 - required answer / reaction
- abilities to generalized to unseen tasks
- Examples:
 - Existing NLP tasks: summarization, translation, text classification, question answering
 - augmenting them with human labeled task descriptions
 - Dialogue system / Open Question Answering
 - human labelers compose instructions for tasks and another group of labelers answer them

Adaptation: Alignment Tuning

Alignment Tuning:

- LLMs sometimes exhibit unintended behavior
 - false information
 - inaccurate objectives
 - harmful, misleading, biased outputs
- Pre-training objective:
 - lack of human values or preferences
- Alignment Tuning: making LLMs act in line with human expectations
 - Criteria:
 - Helpfulness: attempt to solve the given task in a concise and efficient manner
 - Honesty: accurate content instead of fabricated information (hallucination)
 - Harmlessness: not offensive or discriminatory

Adaptation: Alignment Tuning: Human Feedback

Alignment Tuning:

- **Collecting Human Feedback:**

- high quality human feedback
- Human Labeler Selection:
 - human labelers should have a qualified level of education and excellent proficiency
- Human Feedback Collection:
 - Ranking-based:
 - evaluating model-generated outputs by ranking them
 - easier to compare outputs than generate the best
 - not selecting the best answer(s)
 - more consistent answer (than selecting the best)
 - Question-based:
 - human labelers answer certain questions designed by researchers
 - covering the alignment criteria and additional constraints
 - task-specific questions based on the generated output
 - Rule-based:
 - response preference feedback: comparing the quality of output pairs
 - rule violation feedback: scoring the extent of the violation of the rules

Adaptation: Alignment Tuning: RLHF

Alignment Tuning:

- **Reinforcement Learning with Human Feedback (RLHF):**
 - OpenAI's approach
 - fine-tuning LLMs with the collected human feedback using Reinforcement Learning
 - Components:
 - Pre-trained Language Model
 - Reward model:
 - learning from the human feedback
 - provides the training signal for the Language Model
 - reflecting the human preferences
 - Reinforcement Learning algorithm:
 - Proximal Policy Optimization
 - training the Language Model

Adaptation: Alignment Tuning: RLHF

Alignment Tuning:

- Reinforcement Learning with Human Feedback (RLHF):

- Steps of RLHF:
 - Supervised Fine-Tuning:
 - collecting a supervised dataset containing input prompts and desired outputs
 - fine-tuning the Language Model
 - prompts and outputs can be written by human labelers
 - Reward model training:
 - the Language Model generates a certain number of output texts
 - human labelers annotate the preference for these pairs
 - ranking the generated candidate texts
 - a Reward model is trained to predict the human-preferred output
 - RL Fine-Tuning:
 - aligning: fine-tuning the Language Model by a Reinforcement Learning algorithm
 - the pre-trained Language Model acts as the policy:
 - taking as input a prompt
 - returning an output text
 - action space: the vocabulary
 - state: currently generated token sequence
 - reward: provided by the Reward model
 - Regularization:
 - to avoid diverging significantly from the initial Language Model
 - a penalty term: KL-divergence
 - between the generated results from the current Language Model and the initial Language Model

Adaptation: Alignment Tuning: RLHF

Advantages of RLHF:

- helping with hallucination
- Why RLHF works?
 - Diversity:
 - Supervised Fine-Tuning:
 - the LM's output is expected to match the demonstrated responses
 - demonstration only gives the model positive signals
 - no negative signals
 - RLHF:
 - the model gets negative feedback as well
 - RL allows to show negative signals as well

Prompting

Prompting:

- no parameter update to the LLM
- using In-Context Learning abilities:
 - Instructions:
 - specifying what to do
 - Descriptions:
 - detailing and explaining the task
 - Examples:
 - providing example solutions

Comparisons

Pre-training vs. Fine-tuning

Fine-tuning vs. Prompting

Pre-Training vs. Supervised Fine-Tuning

Pre-training stage:

- is about knowledge
- Data:
 - raw text
 - low quality
 - large quantity

Supervised fine-tuning stage:

- is about alignment
- learning to respond more consistently
- increasing knowledge of new specific concepts
- correcting old incorrect information
- Data:
 - demonstration data: in the proper format (prompt, response)
 - high quality data
 - smaller quantity
 - 10,000 - 100,000 pairs
 - format: (prompt, response)
- fine-tuning stage is much cheaper:
- companies do it a lot more frequently than pre-training

Pre-Training vs. Supervised Fine-Tuning

Pre-training stage:

- is about knowledge
- Data:
 - raw text
 - low quality
 - large quantity

Supervised fine-tuning stage:

- is about alignment
- learning to respond more consistently
- increasing knowledge of new specific concepts
- correcting old incorrect information
- Data:
 - demonstration data: in the proper format (prompt, response)
 - high quality data
 - smaller quantity
 - 10,000 - 100,000 pairs
 - format: (prompt, response)
- fine-tuning stage is much cheaper:
- companies do it a lot more frequently than pre-training

Fine-tuning vs. Prompting

Prompting:

- rather go by prompting first then fine-tuning
- Pros:
 - no data to get started
 - smaller upfront cost
 - no technical knowledge needed
 - connect data through retrieval (RAG)
 - lower cost during training
- Cons:
 - much less data fits
 - forgets data
- great for generic, side projects, prototypes

Fine-tuning:

- customizing the model to a specific use case
- easier to show than tell
- typical clear improvements from 50-100 examples
- fine-tuning lets us put more data into the model than what fits into the prompt
- Pros:
 - nearly unlimited data fits
 - not forgetting data
 - reducing hallucination
 - lower cost during inference
- Cons:
 - more high-quality data
 - upfront compute cost
 - needs some technical knowledge, esp. data
- great for domain-specific, enterprise, production usages, ... privacy!

Fine-tuning vs. Prompting

Prompting:

- rather go by prompting first then fine-tuning
- Pros:
 - no data to get started
 - smaller upfront cost
 - no technical knowledge needed
 - connect data through retrieval (RAG)
 - lower cost during training
- Cons:
 - much less data fits
 - forgets data
- great for generic, side projects, prototypes

Fine-tuning:

- customizing the model to a specific use case
- easier to show than tell
- typical clear improvements from 50-100 examples
- fine-tuning lets us put more data into the model than what fits into the prompt
- Pros:
 - nearly unlimited data fits
 - not forgetting data
 - reducing hallucination
 - lower cost during inference
- Cons:
 - more high-quality data
 - upfront compute cost
 - needs some technical knowledge, esp. data
- great for domain-specific, enterprise, production usages, ... privacy!

GPT series

Technical evolution of the GPT-series models:

- 2 key points to the success:
 - training decoder-only Transformer LMs that can accurately predict the next word
 - scaling up the size of LMs
- Initial GPT models:
 - **GPT-1**
 - **GPT-2**
- Large Language Models (LLMs):
 - **GPT-3**
 - **Fine-tuned models**
 - **GPT-3.5**
 - **ChatGPT**
 - **GPT-4**

GPT series

Technical evolution of the GPT-series models:

- 2 key points to the success:
 - training decoder-only Transformer LMs that can accurately predict the next word
 - scaling up the size of LMs
 - Initial GPT models:
 - **GPT-1**
 - **GPT-2**
 - Large Language Models (LLMs):
 - **GPT-3**
 - **Fine-tuned models**
 - **GPT-3.5**
 - **ChatGPT**
 - **GPT-4**
- in 2018
 - the Transformer architecture is adopted (change RNNs)
 - unsupervised pre-training and supervised fine-tuning
 - 117 million parameters

GPT series

Technical evolution of the GPT-series models:

- 2 key points to the success:
 - training decoder-only Transformer LMs that can accurately predict the next word
 - scaling up the size of LMs
 - Initial GPT models:
 - **GPT-1**
 - **GPT-2**
 - Large Language Models (LLMs):
 - **GPT-3**
 - **Fine-tuned models**
 - **GPT-3.5**
 - **ChatGPT**
 - **GPT-4**
- “Unsupervised Multitask Learner”
 - in 2019
 - 1.5 billion parameters
 - unsupervised language modeling (no explicit fine-tuning on labeled data)
 - $p(\text{output} \mid \text{input}, \text{task})$
 - conditioned on the input and the task information
 - NLP tasks can be considered as the word prediction problem based on a subset of the world text
 - unsupervised LMs could be capable in solving various tasks
 - GPT-2 paper: “Language Models are Unsupervised Multitask Learners”
 - inferior performance compared with supervised fine-tuning SOTA methods

GPT series

Technical evolution of the GPT-series models:

- 2 key points to the success:
 - training decoder-only Transformer LMs that can accurately predict the next word
 - scaling up the size of LMs
- Initial GPT models:
 - **GPT-1**
 - **GPT-2**
- Large Language Models (LLMs):
 - **GPT-3**
 - **Fine-tuned models**
 - **GPT-3.5**
 - **ChatGPT**
 - **GPT-4**
- in 2020
 - 175 billion parameters
 - introducing: In-Context Learning (ICL)
 - utilizing LLMs in a few-shot or zero-shot manner
 - ICL can instruct a LLMs to understand the tasks in the form of natural language text
 - pre-training and utilization of LLMs converge to the same language modeling paradigm
 - excellent performance in a variety of NL tasks
 - large performance leap
 - scaling law: larger models have significantly stronger ICL abilities

GPT series

Technical evolution of the GPT-series models:

- 2 key points to the success:
 - training decoder-only Transformer LMs that can accurately predict the next word
 - scaling up the size of LMs
- Initial GPT models:
 - **GPT-1**
 - **GPT-2**
- Large Language Models (LLMs):
 - **GPT-3**
 - **Fine-tuned models**
 - **GPT-3.5**
 - **ChatGPT**
 - **GPT-4**
- 2 major approaches to further improve the GPT-3 model
- Training on code data: **Codex**
 - GPT-3: lacks reasoning ability on complex tasks
 - in 2021
 - a GPT model fine-tuned on a large corpus of GitHub code
 - solving complex coding tasks and math problems
 - improving reasoning abilities and chain-of-thought prompting abilities
- Alignment with human preference: **InstructGPT**
 - in 2022
 - improving GPT-3 for human alignment
 - issues of generating harm or toxic content
 - key to the safe deployment of LLMs in practice
 - 3-stage Reinforcement Learning from Human Feedback (RLHF)

GPT series

Technical evolution of the GPT-series models:

- 2 key points to the success:
 - training decoder-only Transformer LMs that can accurately predict the next word
 - scaling up the size of LMs
- Initial GPT models:
 - **GPT-1**
 - **GPT-2**
- Large Language Models (LLMs):
 - **GPT-3**
 - Fine-tuned models
 - **GPT-3.5**
 - ChatGPT
 - **GPT-4**
- based on a code-based and RLHF-based GPT models
- in 2022

GPT series

Technical evolution of the GPT-series models:

- 2 key points to the success:
 - training decoder-only Transformer LMs that can accurately predict the next word
 - scaling up the size of LMs
- Initial GPT models:
 - **GPT-1**
 - **GPT-2**
- Large Language Models (LLMs):
 - **GPT-3**
 - **Fine-tuned models**
 - **GPT-3.5**
 - **ChatGPT**
 - **GPT-4**
- conversation model
- in 2022
- based on GPT-3.5 (and later GPT-4)
- trained in a similar way as InstructGPT
 - specially optimized for dialogue
- plugin mechanism:
 - further extending the capabilities of ChatGPT with existing tools or apps

GPT series

Technical evolution of the GPT-series models:

- 2 key points to the success:
 - training decoder-only Transformer LMs that can accurately predict the next word
 - scaling up the size of LMs
- Initial GPT models:
 - **GPT-1**
 - **GPT-2**
- Large Language Models (LLMs):
 - **GPT-3**
 - **Fine-tuned models**
 - **GPT-3.5**
 - **ChatGPT**
 - **GPT-4**
- in 2023
- extending the text input to multimodal signals
- better performance on evaluation tasks than GPT-3.5
- Red Teaming:
 - Goal: reducing the harm or toxic content generation
 - using manual or automated methods to adversarially probe a language model for harmful outputs, and then updating the model to avoid such outputs
- Predictable scaling:
 - accurately predicting the final performance with a small proportion of compute during model training

Proprietary vs. Open-source Models

Comparison between closed and open models:

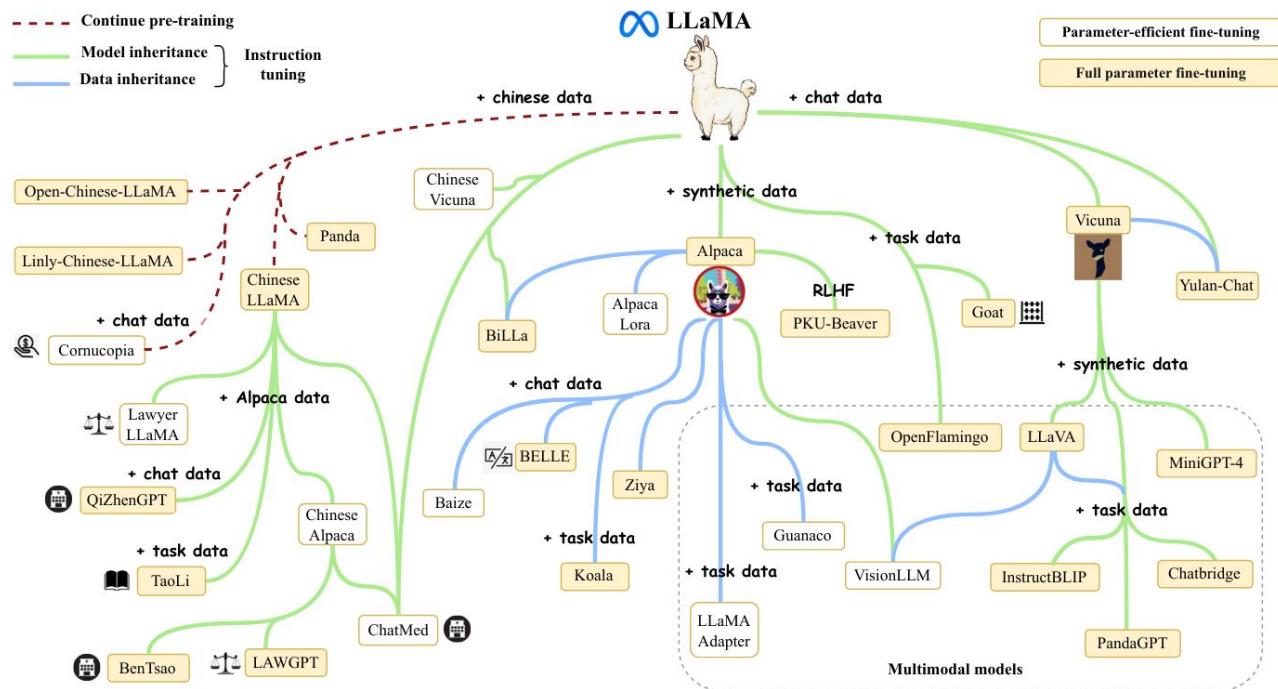
Closed models:

- better performance
- we cannot
 - really work with them
 - fine-tune with them
 - download them

Open-source models:

- lower performance
- depending on our application can be good enough
- we can fine-tune

LLaMa model series



LLaMa model series

LLaMa:

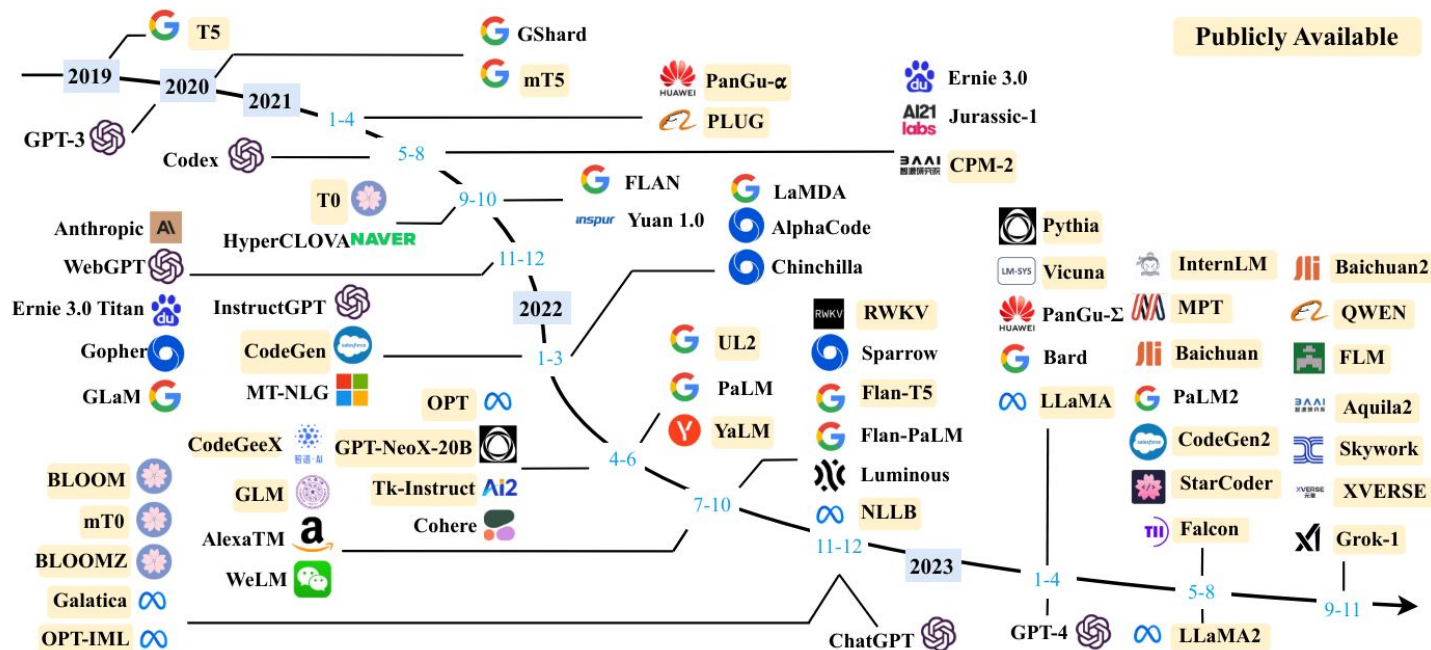
- probably the most powerful open weights model
- weights, architecture, paper: all released by Meta
- example model: llama-2-70b:
- 70 billion parameters:
- a parameters: float 16 → 140 GB of size
- releasing both base model and fine-tuned model
- training:
- data: chunk of the internet: 10TB of text
- 6000 GPUs for 12 days
- 140 GB model size
- 2 million dollars cost

Hungarian LLMs

Magyar NLP modellek:

- https://acta.bibl.u-szeged.hu/78417/1/msznykonf_019_247-262..pdf
- Hugging Face - NYTK modellek (Nyelvtudományi Kutatóközpont):
 - <https://huggingface.co/NYTK>
 - Modellek:
 - PULI-GPTrio
 - named-entity-recognition-nerkor-hubert-hungarian
 - PULI-GPT-3SX
 - sentiment-hts5-xlm-roberta-hungarian
 - sentence-transformers-experimental-hubert-hungarian
 - PULI-GPT-2
 - PULI-BERT-Large
 - Adathalmazok:
 - HuCOLA
 - HuSST
 - HuCoPA
 - HuWNLI
 - HuRC

List of LLMs



List of LLMs

(Large) Language Models:

- BERT: <https://arxiv.org/abs/1810.04805>
- GPT-1: <https://openai.com/research/language-unsupervised>
- GPT-2: <https://openai.com/research/better-language-models>
- GPT-3: <https://openai.com/research/language-models-are-few-shot-learners>
- GPT-4: <https://openai.com/research/gpt-4>
- ChatGPT: <https://openai.com/blog/chatgpt>
- XLNet: <https://arxiv.org/abs/1906.08237>
- RoBERTa: <https://arxiv.org/abs/1907.11692>
- ALBERT: <https://arxiv.org/abs/1909.11942>
- T5: <https://arxiv.org/abs/1910.10683v4>
- ELECTRA: <https://arxiv.org/abs/2003.10555>
- DeBERTa: <https://arxiv.org/abs/2006.03654>
- PaLM: <https://arxiv.org/abs/2204.02311>
- ELMO: <https://arxiv.org/abs/1802.05365>
- ULMFiT: <https://arxiv.org/abs/1801.06146>
- DistilBERT: <https://arxiv.org/abs/1910.01108>
- XLM-RoBERTa: <https://arxiv.org/abs/1911.02116>
- UniLM: <https://arxiv.org/abs/1905.03197>
- StructBERT: <https://arxiv.org/abs/1908.04577>
- MobileBERT: <https://arxiv.org/abs/2004.02984>
- CTRL: <https://arxiv.org/abs/1909.05858>
- Flair: <https://aclanthology.org/N19-4010.pdf>
- Llama: <https://arxiv.org/abs/2302.13971>
- Llama-2: <https://arxiv.org/abs/2307.09288>
- Transformer-XL: <https://arxiv.org/abs/1901.02860>
- ERNIE: <https://arxiv.org/abs/1905.07129>
- BLIP: <https://arxiv.org/abs/2201.12086>
- BLIP-2: <https://arxiv.org/abs/2301.12597>
- Flamingo: <https://arxiv.org/abs/2204.14198>
- OpenFlamingo: <https://arxiv.org/abs/2308.01390>
- LLaVA: <https://llava-vl.github.io/>

ChatGPT API

<https://openai.com/index/introducing-chatgpt-and-whisper-apis>

<https://help.openai.com/en/collections/3675931-api>

Additional resources

[A Survey of Large Language Models](#)